MeshTac: Mesh-Aligned Pseudo-Labels for Label-Efficient Tactile Depth, Normal and Pose Estimation

Extended Version

Anonymous submission

***Abstract*— Known-object tactile perception often has access to object meshes but few real contacts registered to them. Simulating additional images avoids this registration step, yet a learned renderer may itself require substantial real supervision. We present MeshTac, a method that uses a small set of labeled contacts to align known meshes with unlabeled tactile images. A teacher predicts an initial contact state, and a residual aligner trained on perturbations of the labeled contacts refines it. The aligned mesh then supplies depth, normal and pose targets for a student trained on real images. All components share the same real training labels, and the trained student requires no mesh input. Experiments on ridge-pattern objects, with models trained separately for each sensor, show improved contact reconstruction over real-only training, synthetic-data augmentation and teacher-based self-training. These gains persist without a learned renderer, indicating that mesh-pose accuracy is a key bottleneck in recovering supervision from unlabeled contacts. A complementary label-accounting study shows that renderer supervision can substantially change the apparent benefit of simulation. Synthetic data remains useful across capture sessions, and the same aligner can optionally refine poses at inference when the mesh is available. Together, the results establish mesh alignment as a practical route to label-efficient tactile perception for known objects.**

I. Introduction

Recovering contact geometry from touch gives a robot information about where an object meets its fingertip and how the object is oriented. Vision-based tactile sensors obtain this information by imaging the deformation of a soft elastomer. GelSight, GelSlim and DIGIT [1]–[3] follow this principle, but their images differ with illumination, elastomer properties and camera response. Learning a geometric predictor requires examples of the target sensor's appearance together with accurate labels. We consider known-object perception: object identity and a mesh are available, but registering every captured contact to that mesh is costly.

Simulation can reduce the need for registered real contacts by generating tactile images and geometric labels from known meshes [4]–[6]. Its usefulness depends on whether the rendered appearance supports learning on the real sensor. Learned renderers can improve this correspondence, but training them also consumes real labels. A fair assessment of label efficiency must therefore count the supervision used by both the renderer and the downstream model.

Our renderer study exposes the importance of this accounting. A model augmented with images from a renderer trained on all real training pairs appears substantially more accurate than a model trained on a small labeled subset. Much of this advantage disappears when both components are restricted to the same subset: in the few-label regime on sensor A, learned rendering slightly lowers all-pixel depth error while worsening contact error and overlap. This finding motivates a more direct use of the mesh to supervise real images.

An unlabeled tactile image already contains the target sensor's appearance. If the object mesh is known, the missing information for a geometric label is its contact state: in-plane rotation, translation and indentation. A model trained on a small labeled set can initialize this state. Its estimate is often too imprecise to use directly, however, because pose errors shift the projected contact and corrupt the resulting dense labels. The central challenge is to align the mesh accurately enough for the additional images to become useful training data.

MeshTac addresses this challenge with a learned residual aligner (Fig. 1). Given a real image and mesh depth at a candidate state, the aligner predicts a correction. Perturbing the registered contacts supplies training examples without additional real labels. Unlike refinement used only at inference, MeshTac converts each corrected state into depth, normal and pose targets for training on real images. The student retains real contact appearance during training and predicts without a mesh at standard inference.

We evaluate MeshTac against real-only training, synthetic-data augmentation and teacher-based self-training under matched training-label budgets. Experiments with models trained separately for each sensor show that alignment improves both pseudo-label quality and the student's contact reconstruction. Renderer-free and second-session experiments distinguish the contribution of geometric supervision from that of synthetic appearance. Our contributions are:

- **Mesh-aligned pseudo-labeling for known-object tactile perception.** MeshTac learns contact-state refinement from a small registered set and converts the refined meshes into consistent depth, normal and pose targets for unlabeled real images; the trained student remains mesh-free at standard inference.
- **Label accounting for tactile simulation.** A controlled study gives the renderer and downstream model the same real training-label budget, separating gains from synthetic data from gains obtained with additional paired real contacts.
- **Systematic empirical characterization.** Experiments with separately trained models on two sensors, a second capture session and held-out objects isolate the roles of alignment and synthetic data; they also evaluate reuse of

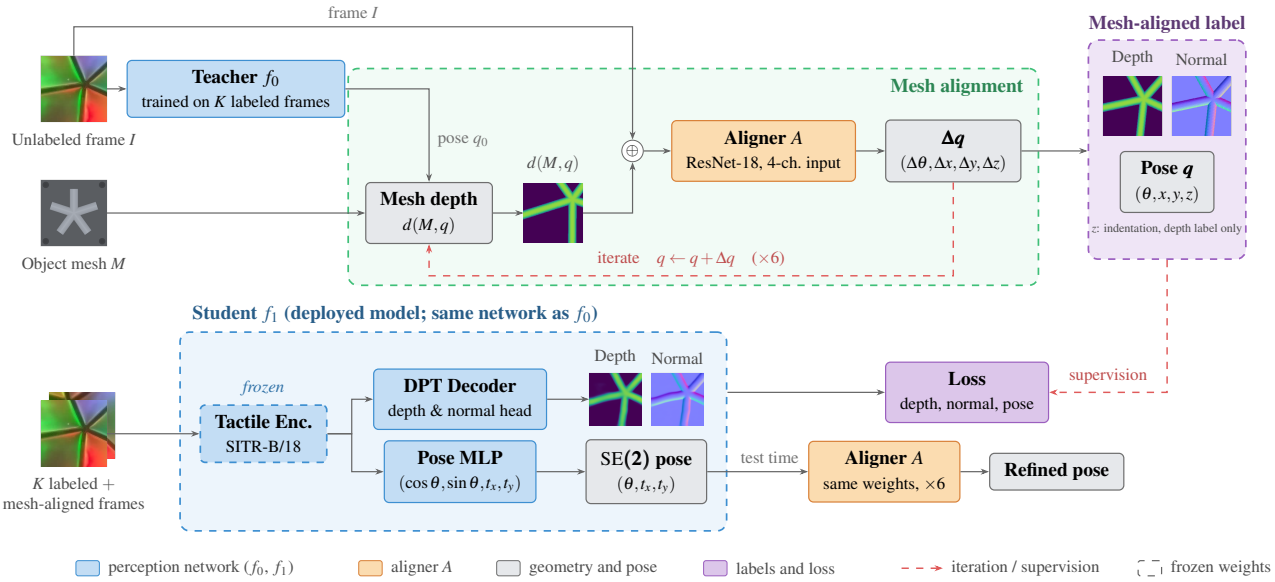


Fig. 1: **Overview of MeshTac.** Teacher f_0 initializes the contact state. Aligner A compares the frame with mesh depth and predicts a correction, applied six times. The aligned mesh supplies depth, normal and planar-pose targets for student f_1, which shares the teacher's architecture. All components use the same K real training labels; the main sensor-A pipeline also uses synthetic examples, omitted here. The student requires no mesh at inference; optional pose refinement uses the image, mesh and fixed aligner. Blue: perception networks; orange: aligner; grey: geometry and pose; purple: labels and loss. Thumbnails show a $K = 25$ example, including the student's actual prediction.

the same aligner for optional test-time pose refinement without retraining.

II. Related Work

Tactile simulation and transfer. Tactile simulation must model both contact geometry and its appearance through the sensor. TACTO [4] emphasizes optical rendering, Taxim [5] models optical response and marker displacement from calibration examples, and Tacchi [6] simulates elastomer deformation. Physics-based [7], differentiable [8] and large-scale simulators [9] offer further choices in contact modeling and computational cost. Learned image translation connects simulated and real tactile observations [10]–[12]; unpaired translation [13]–[15] and domain randomization [16] reduce reliance on paired examples. Our study examines how the real supervision used to fit a renderer affects its downstream benefit. Representation learning addresses sensor variation at the feature level [17]–[19]. We hold this component fixed by using the same frozen SITR encoder [17] throughout.

Pseudo-labeling and pose refinement. Self-training extends supervision to unlabeled data using model predictions [20], [21], but errors in those predictions can carry into student training. Known geometry offers a constraint on the labels. In visual pose estimation, DeepIM [22] learns iterative corrections from observed and rendered images, while Self6D [23] uses unlabeled real images for self-supervised pose learning. MeshTac brings residual refinement to tactile contact states: it compares tactile RGB with candidate mesh depth and learns corrections from perturbed labeled contacts. Its primary use of refinement is label generation rather than inference alone: the corrected mesh supplies dense targets and pose labels for a student that later predicts from tactile images without mesh input. Because alignment compares the observation with mesh depth, this labeling step requires neither a photorealistic renderer nor differentiation through rendering.

III. Data, Models and Evaluation Protocol

A. Problem Statement

We study how to learn tactile depth, normals and pose when object meshes are known but registered real contacts are scarce. The comparison includes synthetic-data augmentation and mesh-aligned pseudo-labeling, with real training labels shared across components.

A tactile renderer maps a mesh M and contact state q (planar pose and indentation) to an image and dense labels, $(\hat{I}, D, N) = \mathcal{R}_\theta(M, q)$. The geometry stage $(M, q) \mapsto (D, N)$ is shared by every renderer we study; renderers differ only in the appearance operator $\mathcal{P}_\theta : (D, N) \mapsto \hat{I}$ and in the real data $\mathcal{D}^r_{\text{cal}}$ used to fit θ. A downstream model f_ϕ predicts depth, normals and planar pose from a real image. We measure a renderer by the real test error of f_ϕ trained on its synthetic data plus a real subset $\mathcal{D}^r_{\text{task}}$, and we control the total real budget $\mathcal{D}^r_{\text{cal}} \cup \mathcal{D}^r_{\text{task}}$.

B. Sensors and Data

Sensor A. sensor A is a GelSlim-type fingertip sensor with a 17.5 mm square gel view imaged under six colored LED emitters. Images are center-cropped to 0.816 of the view and resized to 224×224.

Real data (session 1). A KUKA LBR Med arm presses the sensor onto 12 ridge-pattern objects of known geometry, recording the image and end-effector pose at every contact; ground-truth (GT) depth and normals are the mesh projected at the recorded pose (geometric targets, not membrane

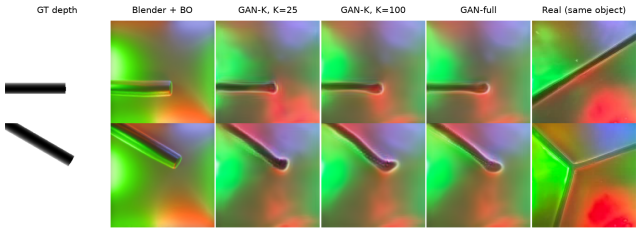


Fig. 2: **Two renderers, one contact.** Each row is one simulated contact (GT depth, left) rendered by the physically calibrated Blender renderer, by the learned renderer trained on $K = 25$ or $K = 100$ real pairs per object, and by the learned renderer trained on all real pairs (GAN-full). Right: a real frame of the same object. The physical renderer reproduces the illumination field but not the dark contact shading of the real gel; the learned renderers reproduce the shading and inherit the illumination of their training frames.

measurements). After removing frames whose image contact region disagrees with the GT, 4,669 frames remain (222–568 per object), one session per object, every tenth frame held out as validation (472 frames).

Real data (session 2). The same objects were re-placed and re-captured on another day (48 runs). GT poses were re-estimated per frame by maximizing the IoU between the image ridge mask and the projected mesh contact and manually reviewed; the curated test set has 854 frames of 11 objects (Sec. V-C). Nothing from this session is used for training or model selection.

Sensor B. GelSlim 4.0 [24] is the previous generation of the same fingertip family, with a marker array, a different illumination layout and a 12 mm field of view. Its public depth dataset [24] has the same 12 objects pressed by a parallel gripper at 1 mm indentation with mesh-derived depth and in-hand pose labels (135 test frames per object). We use the left finger, the central 12×12 mm, and report rotation after removing a per-object constant offset because its pose origin differs from ours.

The public depth labels are clipped at the nominal indentation and smoothed, so they are not exact mesh projections. For the MeshTac comparisons in Sec. V-E, a global search fits the planar mesh pose to each public label (0.018 mm mean residual). We canonicalize the fitted pose using the object's rotational symmetry and use its projected mesh depth for training and testing (*mesh-registered* labels). A separate set of renderer experiments in Sec. V-E retains the original public labels; those results are identified explicitly.

Contact sampler. For each object we sample contact cells over the mesh at 12 base rotations (30° apart), an indentation of 0.8–1.2 mm and in-plane jitter, and export D from the object z-buffer and N by finite differences. Simulated poses are restricted to each object's real rotation and translation ranges. The same 5,345 simulated contacts feed every renderer.

C. Physically Calibrated Renderer

The physical renderer is implemented in Blender/Cycles: a gel surface mesh is shrink-wrapped onto the object and smoothed (35 corrective iterations), shaded with a glossy/transparent material and a height-driven contact-darkening term, lit by six area emitters, and post-processed (saturation, brightness, contrast, haze, a low-frequency random gel warp). Emitter and post-processing parameters (16 scalars) are fitted by tree-based Bayesian optimization of an L_2 plus histogram distance between the real and rendered no-contact image (500 evaluations); gel and contact parameters were hand-tuned on a few contact frames. Thus, one no-contact image is the automatic calibration target, but additional contact images informed manual tuning. Their use must be distinguished from the counted paired training labels.

Domain-randomization baseline. To test whether careful calibration is necessary, we also render with randomized sensor parameters: random HSV colors and strengths for each LED emitter, random gel optics, per session, with the geometry stage (scale, field of view, smoothness) fixed. This variant uses no image to fit its randomized appearance parameters; it retains the shared geometry configuration.

D. Learned Renderer

The learned renderer replaces $\mathcal{P}_\theta$ by a conditional GAN [25]: a five-level U-Net that takes the depth map and two coordinate channels and outputs an RGB image, a 70×70 spectral-normalized PatchGAN discriminator, and an adversarial $+\,100\,(\mathrm{L1}+\mathrm{gradient\ L1}+\mathrm{low\text{-}frequency\ L1})$ loss. Simulated depth is jittered by $\pm 15\%$ at generation time.

Real budget. GAN-K is trained on K real (depth, image) pairs per object, an evenly spaced subset of each object's training frames that is *the same* K frames the downstream model receives; GAN-LOO on all training frames of the 9 non-held-out objects; GAN-full on all 4,197 pairs. Small budgets train for at least 30k generator updates.

Variants. *Background-conditioned*: the session's no-contact image is appended as three input channels, providing session illumination from one background frame. *Physics-conditioned*: the physical renderer's image of the same contact is appended, optionally predicting only a residual on top of it, optionally with the background. *Difference image*: the target is the image minus the session background. *Blender-pretrained*: pretraining on 30k (depth, Blender image) pairs, then fine-tuning on the real pairs.

Multi-sensor: the background-conditioned GAN is trained on 100 simulated sensor designs (random LED layouts and gel materials) plus optionally the target sensor's real pairs; this tests whether sensor diversity can replace sensor-specific real data. *Rotation coverage* changes the geometry stage instead: we vary how many of an object's real rotations appear in simulation (1, 2 or all 12) and, for the learned renderer, rotate the depth map before rendering so that every rotation is covered.

E. Zero-Label Alternative: Contact-Preserving Unpaired Refiner

We also test whether unlabeled real images can improve synthetic appearance without paired renderer supervision. Following unpaired image translation [11], [13], [15], we train a refiner to change the shading of a physical render

while preserving its contact region. The refiner $\mathcal{F}$ takes the physical render $\hat{I}_{\text{phys}}$, the session background B and the contact mask m (simulated depth > 0, dilated by 4 px and feathered) and outputs

$$\hat{I} = m \odot \mathcal{F}(\hat{I}_{\text{phys}}, B, m) + (1 - m) \odot B, \tag{1}$$

The background is therefore taken from the real sensor, while the refiner changes only the contact region. Depth, normal and pose targets remain those of the physical render.

$\mathcal{F}$ is the same U-Net as the learned renderer with a 7-channel input. It is trained against a 70×70 PatchGAN on *unlabeled* real frames (image only; no depth, no pose) with an adversarial loss and a patchwise contrastive loss (PatchNCE [15]) between the encoder features of input and output, sampled preferentially inside m, to encourage preservation of the rendered ridge structure. We train for 10k generator updates with batch size 8 and Adam (2×10^{-4}, $\beta_1 = 0.5$).

Its real budget is the set of unlabeled frames plus one background frame per session, and the same leakage rule applies: in leave-object-out the held-out objects' frames are excluded, in K-shot the same K frames as the downstream model are used without their labels. With domain-randomized inputs, the refiner requires no appearance calibration or paired training labels; it retains the shared geometry configuration.

F. Downstream Tactile Model

The downstream model is identical for every renderer: a frozen sensor-invariant encoder (SITR-B/18 [17]) provides patch tokens at four depths; a DPT decoder [26] predicts depth and normals at 224×224; a pose head regresses $(\cos\theta, \sin\theta, t_x, t_y)$ from the class token and mean patch token. For the classification study, we add a 12-class head (LayerNorm on concatenated class and mean patch tokens, two-layer MLP with GELU) and train with cross-entropy alongside the other losses. The decoder, pose head and object embedding are trained (12M parameters). Losses are MSE on depth and normals, $1 - \cos(\theta - \hat{\theta})$ on rotation and L1 on translation, balanced by grouped uncertainty weighting [27].

When synthetic data is used, training mixes simulated and real frames. Simulated frames receive photometric augmentation and $\pm 15^\circ$ gel-spin rotation with matched labels; real frames are not augmented. We use AdamW with learning rate 2×10^{-4}, batch size 32 and a cosine schedule. Main sensor-A runs use 150 epochs, while sensor B and seed/ablation studies use 50. Depth and normal metrics use the checkpoint with the lowest validation depth MSE, sensor-A pose metrics in the main within-session and pose-refinement tables use the one with the lowest validation angular error. The legacy cross-session table reports pose from the depth checkpoint. In every case, checkpoint selection uses validation labels in addition to the K training labels.

G. Data-Access Protocol

Paired renderers and the aligner exclude validation and test images from training. The K-shot protocol additionally limits paired training labels to the same K frames per object. Validation labels remain available for model selection; the reported budget counts training labels. Physical-renderer manual tuning is described separately above.

Leave-object-out (LOO). Three objects (a two-line, a rod and a cross pattern; 1,244 frames) are removed from all real training, excluded from paired-renderer and aligner training, and evaluated in full. Simulated contacts of the three meshes are available to every renderer ("new object with a CAD model but no real touches"). A *geometry control* additionally removes the three meshes from simulation.

K-shot. Every object contributes $K \in \{0, 10, 25, 50, 100\}$ evenly spaced real training frames. For $K > 0$, the learned renderer is trained on exactly those frames; the downstream model co-trains on them plus simulated data, or on them alone with the same number of updates. Evaluation is on the real validation split. The aligner of MeshTac receives the same K frames, and the remaining training frames may be used *without* their labels.

Cross-session and cross-sensor. Session-2 and sensor-B test frames are never used to train or select anything in the sensor-A study. In the sensor-B study the K-shot protocol is repeated with sensor-B real pairs and the sensor-B test split.

Leaky reference. GAN-full is trained on all 4,197 sensor-A training pairs. In LOO, these include real pairs from held-out objects; in K-shot, they include training labels withheld from budget-matched methods, but no validation frames. We retain the label "leaky renderer" for this non-comparable reference, shown in italics and excluded from best-result highlighting.

Metrics. We report depth MSE over all pixels (mm^2), following prior work on this sensor. Because $\approx 90\%$ of pixels are background, we also report mean absolute error inside the GT contact region (contact MAE, mm) and contact IoU (predicted depth > 0.05 mm vs. GT > 0), which also penalizes false contact. Normals are evaluated by mean angular error inside the contact region (deg), and pose by mean rotation error (deg) and normalized translation L1.

IV. MeshTac: Mesh-Aligned Pseudo-Labeling

MeshTac converts a small set of registered contacts and a larger pool of unlabeled real images into training data for a tactile predictor. The method first estimates each unlabeled contact state, then projects the known mesh to obtain mutually consistent depth, normal and pose targets.

Geometric labels. Let $d(M, q)$ denote the depth of object mesh M at contact state $q = (\theta, x, y, z)$, where z is indentation. For a known object, these variables determine the mesh-derived depth map, its normals and the planar pose. Refinement therefore takes place in contact-state space, with the object geometry held fixed.

Learning the aligner. A ResNet-18, initialized from ImageNet, takes tactile RGB and candidate mesh depth as four input channels. It predicts a residual A in contact-state units:

$$A\big(I, d(M, q^* + \delta)\big) \approx -\delta, \tag{2}$$

where q^* is the labeled state and δ a sampled perturbation. The ranges are $\pm 8^\circ$, ± 1.5 mm in-plane and ± 0.15 mm in indentation; half the samples use $0.3\times$ these ranges. Residuals

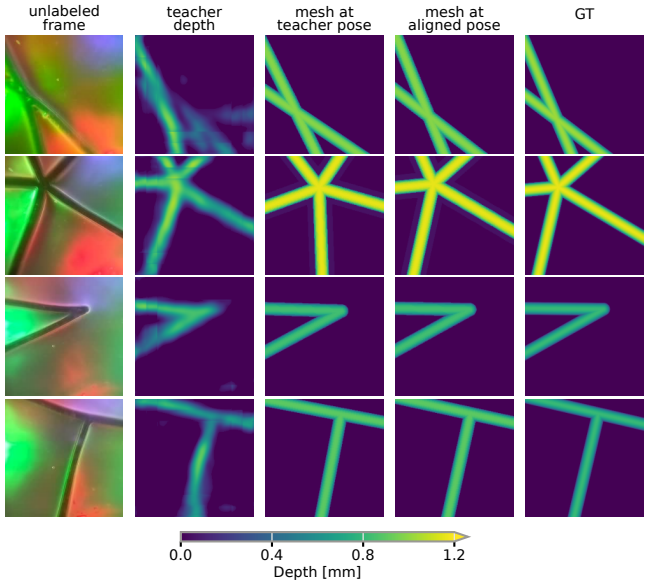


Fig. 3: **Pseudo-label construction.** Sensor A, $K = 25$. Columns compare teacher depth and mesh projections before and after alignment with GT. Each row is a median-IoU example from a subsampled candidate set for one object. All depth maps use the same scale (mm).

are normalized by these ranges for a smooth-L1 training loss (8,000 steps, batch 64, AdamW, learning rate 3×10^{-4}). The main sensor-A aligner mixes 24 real and 40 synthetic examples per batch, with synthetic images generated by the K-pair renderer. The renderer-free and sensor-B aligners use real examples only. Each configuration uses only the same K real training labels available to its teacher.

Labeling real images. For each unlabeled training image, teacher f_0 initializes planar pose and sets indentation to its maximum predicted depth. On sensor A the teacher uses the K real frames and learned-renderer data; the renderer-free variant and sensor B use a real-only teacher. Given the known mesh, we perform six updates,

$$q_{t+1} = q_t + A\big(I, d(M, q_t)\big), \tag{3}$$

clipping indentation to $[0.3, 1.6]$ mm. The final mesh depth, normals and planar pose label the image. We retain every pseudo-labeled training frame in the evaluated configuration (Sec. V-B).

Training the student. Student f_1 has the teacher's architecture and learns from the K labeled frames and the mesh-aligned frames. The sensor-A configuration also uses synthetic images from a renderer retrained on labeled and aligned real pairs. We evaluate a version with no renderer at any stage to isolate alignment. Pseudo-labels are generated exclusively from the training split.

Cost. On one RTX 6000 Ada GPU, training the aligner takes about 20 minutes, labeling 0.5 s per frame (30 minutes for the 3,897 sensor-A frames), and training the student 1.3 hours for 50 epochs or 2.2 hours for 150; the renderer, when used, adds about one hour.

Optional test-time refinement. The student predicts from an image without mesh rendering. When the object identity and mesh are available, the same fixed aligner can refine its predicted pose using the test image. We evaluate this additional inference step separately in Sec. V-D.

V. Experiments

The experiments examine label efficiency, the contribution of alignment, and performance under changes in session or sensor. We then study renderer supervision and appearance variants to explain when synthetic data helps.

Sensor-A main runs use 150 epochs; seed checks, ablations and sensor-B runs use 50 (other exceptions are noted). Metrics are evaluated on real images; depth MSE in mm^2, contact MAE in mm, normals and rotation in degrees. Sections V-A–V-F evaluate MeshTac; Sections V-G–V-I report the renderer comparisons under explicit data-access constraints.

A. Label Efficiency

Tables I and II, and Fig. 4 compare, under the same training-label budget, real-only training, co-training with either renderer, self-training on the teacher's own predictions (Sec. V-B) and MeshTac. Each MeshTac run uses the 2,997–4,077 training frames outside the K budget without their labels.

Simulation alone. Neither paired-renderer baseline improves both contact MAE and IoU for $K \leq 25$; the learned renderer helps from $K = 50$ (depth MSE -10%, -16% at $K = 100$), the physical renderer by at most 5%. Both lower rotation error at $K = 10$ ($9.6° \rightarrow 5$–$6°$) and raise it at $K = 100$.

MeshTac. Mesh-aligned supervision improves dense prediction at every evaluated budget: depth MSE by 13%, 19%, 36% and 37% and contact MAE by 13%, 20%, 27% and 28% at $K = 10$, 25, 50 and 100, with IoU $+0.06$ to $+0.08$ and normal error $-2.6°$ to $-5.1°$. Over three seeds (50 epochs) the depth MSE is 0.0558 ± 0.0003, 0.0447 ± 0.0009, 0.0313 ± 0.0007 and 0.0229 ± 0.0003 against 0.0635 ± 0.0007, 0.0551 ± 0.0010, 0.0480 ± 0.0006 and 0.0366 ± 0.0008 for real-only training.

The improvement also reduces the number of labels needed to reach a given depth error. MeshTac with 50 labels per object (0.0307) has lower depth MSE than real-only training with 100 (0.0369) and is comparable to the learned renderer at 100 (0.0311). At $K = 100$ it has lower error than the all-pairs renderer (depth MSE 0.0232 vs. 0.0274, contact MAE 0.168 vs. 0.205), and at $K = 25$ it has lower contact MAE (0.249 vs. 0.262) and slightly higher IoU, despite 8% higher depth MSE.

The pose results are less uniform. MeshTac has the lowest budget-matched rotation error for $K \leq 50$ ($4.0°$, $3.7°$, $3.7°$); at $K = 100$ real-only training is ahead ($3.2°$ vs. $3.7°$) until the aligner is also used at test time (Sec. V-D).

The depth gains persist without a renderer: with a real-only teacher, an aligner trained on the K frames alone and no synthetic frame at all ("no renderer"), depth MSE is 0.0521, 0.0433, 0.0291 and 0.0227 and rotation error 2.5–4.7°, with lower depth MSE than the renderer-assisted configuration;

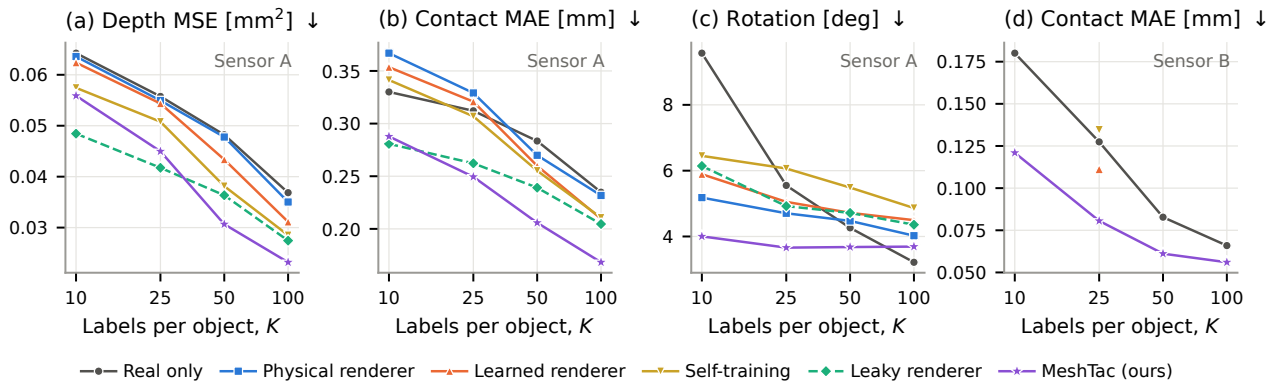


Fig. 4: **Label efficiency.** (a–c) Sensor-A validation results; (d) sensor-B test results with renderer-free MeshTac. Lower is better. Lines connect evaluated budgets. Sensor A uses 150 epochs except the $K = 10$ leaky reference (50); sensor B uses 50.

K	Method	Depth MSE↓	Contact MAE↓	IoU ↑	Normal (deg)↓	Rot. (deg)↓
0	Blender only	**0.0793**	**0.408**	**0.572**	**48.6**	**8.4**
	GAN-full (leaky) only	*0.0459*	*0.274*	*0.705*	*44.6*	*7.0*
10	Real only	0.0642	0.330	0.625	29.8	9.6
	Blender	0.0636	0.367	0.602	30.8	5.2
	GAN-10	0.0624	0.354	0.597	30.6	5.9
	Teacher-label self-training	0.0574	0.341	0.592	28.3	6.5
	MeshTac, L1 search	0.0545	0.320	0.639	27.3	4.6
	MeshTac (with renderer)	0.0559	**0.288**	**0.682**	27.2	4.0
	no synthetic frames	0.0571	0.293	0.678	27.3	**3.2**
	MeshTac (no renderer)	**0.0521**	0.293	0.660	**26.5**	4.7
	GAN-full (leaky)	*0.0484*	*0.281*	*0.693*	*28.5*	*6.1*
25	Real only	0.0558	0.312	0.664	28.2	5.6
	Blender	0.0549	0.329	0.656	27.5	4.7
	GAN-25	0.0544	0.321	0.626	28.3	5.1
	Blender ⊕ GAN-25	0.0554	0.324	0.667	27.4	4.8
	Refiner-25 (0 labels)	0.0551	0.297	0.672	27.1	5.5
	Refiner-U (0 labels, all unlabeled)	0.0566	0.309	0.666	27.8	5.6
	Self-trained GAN-25 (renderer only)	0.0538	0.326	0.641	28.0	4.8
	Teacher-label self-training, model only	0.0517	0.300	0.611	26.5	8.2
	center window only (no stitching)	0.0518	0.328	0.643	26.7	8.4
	Teacher-label self-training	0.0508	0.307	0.628	26.5	6.1
	round 2	0.0500	0.310	0.642	26.0	6.2
	MeshTac, L1 search	0.0482	0.291	0.655	25.5	4.5
	MeshTac (with renderer)	0.0450	0.249	**0.727**	24.0	3.7
	no synthetic frames	0.0456	0.259	0.715	24.7	**2.7**
	MeshTac (no renderer)	**0.0433**	**0.245**	0.726	**23.3**	2.8
	GAN-full (leaky)	*0.0417*	*0.262*	*0.724*	*25.5*	*4.9*

TABLE I: **K-shot with a shared real budget.** Every object contributes K real frames; GAN-K is trained on the same frames. MeshTac and self-trained rows additionally use the unlabeled training frames (Sec. V-B); refiner training uses no paired label. Real validation split (472 frames). Leaky rows use the renderer trained on all real pairs (the $K = 10$ leaky row is a 50-epoch run).

the renderer provides an additional benefit across sessions (Sec. V-C).

Seeds. Table III reports the comparison over three seeds. The depth-MSE differences between MeshTac and real-only training exceed the variation across these runs.

B. Why Alignment: Ablation

We use the same unlabeled images to compare teacher predictions, render-and-compare search and learned alignment (Table IV). Evaluating both label quality and student performance connects the supervision changes to their downstream effect.

Self-training. With the teacher's own depth as label (stitched from five windows to cover the full frame, pose loss on the labeled frames only), depth MSE improves by 8% but contact MAE and IoU do not: the pseudo labels inherit the teacher's localization error (label IoU 0.56), indicating that a lower all-pixel error need not imply improved localization. Adding the teacher's pose as a label improves rotation accuracy with little change in dense metrics, and a second round of self-training changes depth MSE by at most 4% in either direction.

L1 render-and-compare. Replacing the aligner by a search—

K	Method	Depth MSE↓	Contact MAE↓	IoU ↑	Normal (deg)↓	Rot. (deg)↓
50	Real only	0.0482	0.283	0.690	25.5	4.3
	Blender	0.0478	0.270	0.701	25.2	4.5
	GAN-50	0.0434	0.260	0.716	24.3	4.7
	Teacher-label self-training	0.0381	0.255	0.706	22.7	5.5
	MeshTac, L1 search	0.0384	0.247	0.714	22.7	4.1
	MeshTac (with renderer)	0.0307	0.206	**0.769**	20.4	3.7
	no synthetic frames	0.0332	0.215	0.764	21.1	3.0
	MeshTac (no renderer)	**0.0291**	**0.197**	0.761	**19.8**	**2.7**
	GAN-full (leaky)	*0.0363*	*0.239*	*0.747*	*22.7*	*4.7*
100	Real only	0.0369	0.235	0.744	22.1	3.2
	Blender	0.0350	0.232	0.727	22.0	4.0
	GAN-100	0.0311	0.210	0.773	20.8	4.5
	Blender ⊕ GAN-100	0.0329	0.213	0.760	20.5	4.2
	Refiner-100 (0 labels)	0.0340	0.224	0.761	21.1	4.9
	Refiner-U (0 labels, all unlabeled)	0.0329	0.228	0.747	21.1	4.9
	Self-trained GAN-100 (renderer only)	0.0308	0.207	0.774	20.6	4.4
	Teacher-label self-training, model only	0.0290	0.221	0.744	19.8	4.5
	Teacher-label self-training	0.0285	0.211	0.750	19.7	4.9
	round 2	0.0295	0.217	0.736	20.2	4.9
	MeshTac, L1 search	0.0276	0.203	0.771	19.4	4.1
	MeshTac (with renderer)	0.0232	0.168	0.803	17.1	3.7
	no synthetic frames	0.0237	0.168	0.809	17.1	2.7
	MeshTac (no renderer)	**0.0227**	**0.166**	**0.818**	**16.9**	**2.5**
	GAN-full (leaky)	*0.0274*	*0.205*	*0.777*	*19.9*	*4.4*

TABLE II: **Larger training-label budgets.** Continuation of Table I, $K = 50$ and 100. All data-access and metric conventions are unchanged.

K	Method	Depth MSE	Cont. MAE	IoU	Rot.
25	Real only	0.0551 ± 0.0010	0.314 ± 0.010	0.665 ± 0.004	5.5 ± 0.0
25	GAN-25	0.0536 ± 0.0013	0.314 ± 0.008	0.655 ± 0.032	6.0 ± 0.1
25	Teacher-label self-training	0.0507 ± 0.0001	0.310 ± 0.004	0.624 ± 0.002	7.2 ± 0.1
25	MeshTac, L1 search	0.0482 ± 0.0004	0.289 ± 0.004	0.648 ± 0.010	5.5 ± 0.1
25	MeshTac (with renderer)	0.0447 ± 0.0009	0.246 ± 0.002	0.729 ± 0.002	4.9 ± 0.1
25	MeshTac	0.0433 ± 0.0004	0.247 ± 0.000	0.722 ± 0.000	4.0 ± 0.0
100	Real only	0.0366 ± 0.0008	0.234 ± 0.006	0.750 ± 0.004	3.9 ± 0.0
100	GAN-100	0.0299 ± 0.0004	0.201 ± 0.002	0.779 ± 0.001	5.5 ± 0.0
100	Teacher-label self-training	0.0280 ± 0.0001	0.209 ± 0.002	0.743 ± 0.007	6.1 ± 0.0
100	MeshTac, L1 search	0.0275 ± 0.0002	0.201 ± 0.003	0.767 ± 0.005	5.0 ± 0.1
100	MeshTac (with renderer)	0.0229 ± 0.0003	0.164 ± 0.000	0.818 ± 0.000	4.8 ± 0.0
100	MeshTac	0.0222 ± 0.0004	0.163 ± 0.002	0.821 ± 0.002	3.8 ± 0.0

TABLE III: **Seeds.** Mean ± standard deviation over three seeds (50 epochs) on the real validation split. "MeshTac": real-only teacher, no synthetic frame; "with renderer": the configuration of Table I.

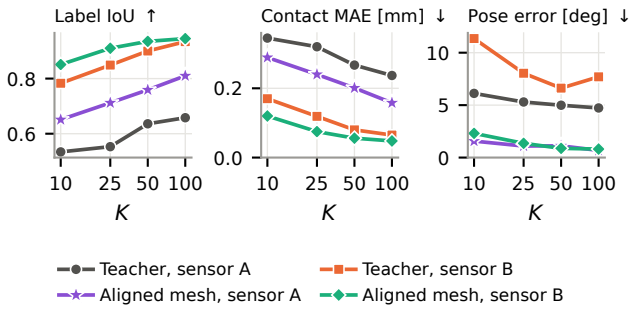


Fig. 5: **Pseudo-label quality.** Teacher predictions and aligned mesh labels, evaluated against withheld training GT on each sensor. The withheld GT is not used for learning.

fit the mesh to the teacher depth on a pose grid, then pick, among 567 neighbouring poses, the one whose GAN-K render is closest to the frame in image L1—helps less (label IoU 0.63 with the half-frame gate, pose error 2.5°) and requires gating: on the half of the frames with the larger render error the mesh label is no better than the teacher's. Image similarity does not reliably identify the correct pose with this few-shot renderer; at $K = 25$ its render at the true pose is farther from the frame than its render at the selected pose on 93% of the frames. This measures score ordering, rather than a pose-failure rate.

Learned aligner. The aligner lowers the teacher's pose error from 5.3° to 1.1° on the 3,897 unlabeled frames and yields labels with IoU 0.71 and contact MAE 0.24 mm (teacher: 0.55, 0.32). Label quality grows with the budget (IoU 0.65, 0.71, 0.76, 0.81 and pose error 1.6°, 1.1°, 1.1°, 0.7° for

Labels of U		Pseudo-label quality			Downstream			
		IoU	Cont.	Rot.	Depth	Cont.	IoU	Rot.
Depth	Pose	↑	MAE↓	(deg)↓	MSE↓	MAE↓	↑	(deg)↓
teacher	–	0.56	0.32	–	0.0507	0.310	0.624	7.2
teacher	teacher	0.56	0.32	4.9	0.0508	0.310	0.622	5.9
mesh$^{1/2}$	L1$^{1/2}$	0.63	0.30	1.4	0.0485	0.286	0.639	6.3
mesh	L1	0.64	0.30	2.5	0.0489	0.292	0.657	5.4
mesh$^{1/2}$	L1	0.63	0.30	2.5	0.0482	0.289	0.648	5.5
mesh	aligner‡	0.71	0.24	1.1	0.0447	0.248	0.730	5.0
mesh	aligner†	0.71	0.24	1.1	0.0447	0.246	0.729	4.9

TABLE IV: **Alignment ablation** ($K = 25$, 50 epochs; first, fifth and last row are three-seed means). Left: quality of the pseudo labels of the unlabeled frames against their withheld GT labels (evaluation only). Right: real validation error of the model trained on them. Depth label: teacher prediction or mesh depth at the fitted pose ($^{1/2}$: mesh on the half of each object's frames with the lowest render error, teacher on the rest). Pose label: none (pose loss on the K labeled frames only), teacher pose, pose found by L1 render-and-compare search with GAN-K ($^{1/2}$: on the confident half only), or the learned aligner. ‡: synthetic data from the teacher-pair renderer; †: MeshTac (renderer retrained on the aligned pairs).

$K = 10$–100). Relative to teacher-depth self-training, the student lowers contact error ($0.310 \rightarrow 0.246$) and raises IoU ($0.624 \rightarrow 0.729$).

Aligner training. The aligner can be trained on real examples alone (label IoU 0.72), while the main mixed real/synthetic configuration gives 0.71; synthetic-only training is less accurate (0.56). Label IoU changes little across the tested settings (half or double the perturbation range: 0.72, 0.73; 2k or 20k steps: 0.73, 0.72; ResNet-34 or -50 backbone: 0.73, 0.71), while dropping the ImageNet initialization costs accuracy (0.67).

Synthetic student data. Student depth MSE is unchanged whether its synthetic frames come from the renderer retrained on aligned pairs or on teacher-labeled pairs (depth MSE 0.0447 in both cases), and within-session depth gains also persist without synthetic data (last row: 0.0433 ± 0.0004 over three seeds with a real-only teacher and no synthetic frame, against 0.0447 ± 0.0009 with the renderer).

Unlabeled data and filtering. Alignment improves labels even on the frames rejected by the render-error gate (IoU 0.66 vs. 0.55 for the teacher), supporting their inclusion in training. A quarter of the unlabeled frames already gives most of the depth-MSE gain (depth MSE 0.0453, 0.0452 and 0.0447 with 25%, 50% and 100% of U), while IoU keeps growing with their number (0.710, 0.725, 0.729).

Repeated alignment. A second round, with the student as the new teacher, offers little further improvement: although the student proposes poses 1.5° closer than the original teacher, the aligned labels (IoU 0.65, 0.71 at $K = 10$, 25) and the retrained student (depth MSE 0.0570, 0.0445; IoU 0.674, 0.717) are within seed variation of the first round.

An unpaired sim-to-real refiner of the physical render, which uses the unlabeled frames without any pose, does not improve on its input (contact MAE 0.349 vs. 0.346 on unseen objects).

Self-trained renderer. Following self-training [20], [21], the K-shot downstream model pseudo-labels U (depth predicted on five windows and stitched back to the full frame; frames kept when the predicted contact overlaps the image's black-hat ridge mask, IoU ≥ 0.2: 3,260 frames at $K = 25$, 2,677 at $K = 100$), and GAN-K is retrained on the K true plus the pseudo pairs. The pseudo labels are as accurate as the labeler (0.32 and 0.24 mm contact MAE against the withheld GT labels), and the renderer's fidelity improves by 5% at both budgets ($0.108 \rightarrow 0.102$, $0.077 \rightarrow 0.073$), while downstream depth MSE improves by about 1%: depth MSE $0.0544 \rightarrow 0.0538$ at $K = 25$ and $0.0311 \rightarrow 0.0308$ at $K = 100$ (contact MAE 0.207, lower than the paired-renderer baseline at $K = 100$). With teacher labels the renderer thus gains little from more real images.

Aligned pairs improve renderer fidelity further: contact-region error falls by 13–19% at every budget ($0.121 \rightarrow 0.105$ at $K = 10$, $0.108 \rightarrow 0.090$ at 25, $0.095 \rightarrow 0.077$ at 50, $0.077 \rightarrow 0.065$ at 100; leaky 0.057), the lowest contact-region errors among the budget-matched renderers. Replacing the teacher-pair renderer with this aligned-pair renderer nevertheless leaves the student's depth MSE unchanged (Table IV). In this within-session comparison, improved synthetic appearance adds little once the real frames carry aligned labels.

Window stitching. Without alignment the label is the teacher depth, and labeling U on the center window alone leaves the pseudo depth empty outside the 0.816 crop (pseudo-label contact MAE 0.41 mm over the full frame instead of 0.32). The self-trained model then has the same depth MSE (0.0518 vs. 0.0517, model only, $K = 25$; Table I) but a worse contact MAE (0.328 vs. 0.300) and a 30% larger peak-depth error, and the gap persists across sessions (contact MAE 0.358 vs. 0.318): the rotation augmentation brings the frame border into the training crop, where unstitched labels are wrong.

Rounds. Without alignment, relabeling U with f_1 and repeating the renderer and model steps improves the pseudo labels at $K = 25$ (label IoU $0.56 \rightarrow 0.60$, contact MAE $0.32 \rightarrow 0.31$ mm, 3,516 frames kept instead of 3,260) but not at $K = 100$ ($0.66 \rightarrow 0.65$), and the downstream change is small in both directions: depth MSE $0.0508 \rightarrow 0.0500$ and IoU $0.628 \rightarrow 0.642$ at $K = 25$, $0.0285 \rightarrow 0.0295$ and $0.750 \rightarrow 0.736$ at $K = 100$, rotation unchanged (6.2°, 4.9°). Additional self-training rounds have limited benefit in these runs. A second MeshTac round likewise yields little further improvement, as reported above.

Selection of the labeled frames. Table VI repeats the renderer-free pipeline with two random selections of the 25 labeled frames per object. Label quality (IoU 0.72, 0.71; pose error 1.8°, 1.6°) and student error are within seed variation of the evenly spaced selection, while real-only training degrades slightly on the random selections.

Where alignment fails. Against the withheld GT of the 3,897 unlabeled frames ($K = 25$), the aligned pose has a mean rotation error of 1.07° and a median of 0.47°; 3.6% of the frames exceed 2°, 1.6% exceed 5° and 1.2% exceed 10° (teacher: 65%, 35%, 16%). The failures concentrate on the five-line object pattern_06 (16.5% of its frames above

Method	Contact MAE↓	IoU ↑	Normal (deg)↓	Rot. (deg)↓	Rot.* (deg)↓
Real only (9 objects, session 1)	0.315	0.675	25.4	16.4	6.5
Blender	0.294	0.701	24.9	11.5	5.1
Blender (domain rand.)	0.294	0.691	25.5	12.3	5.9
Blender, bg-norm.	0.291	0.698	24.9	**9.7**	**4.7**
GAN-LOO	0.272	0.723	24.1	11.1	5.2
Bg-cond. GAN-LOO	0.270	0.727	24.2	10.5	4.7
Blender $\oplus$ GAN-LOO	0.275	0.724	24.0	9.8	4.8
Refiner (0 labels)	0.290	0.714	24.6	12.0	5.3
GAN-full (leaky)	*0.259*	*0.748*	*22.9*	*10.1*	*4.8*
Blender only ($K = 0$)	0.425	0.561	50.7	25.2	5.6
$K = 25$ real only	0.368	0.598	30.7	20.8	5.4
$K = 25$ + Blender	0.382	0.615	30.1	23.0	5.4
$K = 25$ + GAN-25	0.342	0.589	28.7	25.7	5.3
$K = 25$ + refiner-25 (0 labels)	0.338	0.635	29.5	15.5	5.6
$K = 25$ + self-trained GAN-25	0.363	0.606	29.1	20.4	5.5
$K = 25$ + teacher-label self-training	0.335	0.598	27.5	13.1	5.5
$K = 25$ + MeshTac, L1 search	0.322	0.621	27.4	17.2	5.1
$K = 25$ + MeshTac	0.310	0.660	27.2	13.8	5.3
$K = 100$ real only	0.351	0.657	27.5	17.5	5.5
$K = 100$ + Blender	0.299	0.647	25.3	16.4	5.1
$K = 100$ + GAN-100	0.295	0.713	25.6	13.7	5.4
$K = 100$ + refiner-100 (0 labels)	0.330	0.688	26.2	16.9	6.1
$K = 100$ + self-trained GAN-100	0.310	0.715	26.1	12.1	4.9
$K = 100$ + teacher-label self-training	0.300	0.707	24.2	11.0	4.9
$K = 100$ + MeshTac, L1 search	0.293	0.718	23.5	12.5	5.0
$K = 100$ + MeshTac	**0.252**	**0.743**	**22.8**	14.4	5.2

TABLE V: **Cross-session.** Sensor-A models evaluated on the curated session-2 set (854 frames, 11 objects; poses re-estimated per frame and manually reviewed). Rot.: raw; Rot.*: after removing a per-object constant offset (the re-placement of the object). Rows without a K use the session-1 training set of nine objects; object coverage therefore differs from the K-shot rows. Pose metrics use the depth checkpoint.

Training	Depth MSE↓	Contact MAE↓	IoU ↑
Real only, selection 0	0.0558	0.303	0.666
Real only, selection 1	0.0578	0.314	0.655
Real only, selection 2	0.0585	0.324	0.650
MeshTac (no renderer), selection 0	0.0435	0.247	0.722
MeshTac (no renderer), selection 1	0.0426	0.252	0.724
MeshTac (no renderer), selection 2	0.0425	0.244	0.718

TABLE VI: **Robustness to the choice of the K labeled frames** (sensor A, $K = 25$, 50 epochs, real validation split). Selection 0 is the evenly spaced selection used elsewhere; selections 1 and 2 are random permutations of the training frames with the validation split unchanged. Every component (teacher, aligner, student) is retrained on the new selection.

$5°$; every other object below 4%), whose near-symmetric ridge layout admits a second plausible alignment, and the worst 5% of frames carry 53% of the summed rotation error (mean $0.53°$ without them). Translation error (mean 0.98 mm, median 0.52) is dominated by shifts along the ridge direction: 73% of the rod frames and 20–40% of the frames of the line patterns exceed 1 mm, against 2–6% for the objects with crossing ridges (pattern_32, 33, 35, 36). Such shifts leave the depth label unchanged, which is why the label IoU is high despite them. At $K = 100$ the mean rotation error is $0.68°$ and 1.3% of the frames exceed $5°$.

C. A Second Session

The second session tests changes in appearance and object placement without retraining (Table V). MeshTac reduces contact MAE relative to the budget-matched real-only models by 16% at $K = 25$ and 28% at $K = 100$, with MAE/IoU of $0.310/0.660$ and $0.252/0.743$, respectively. The nine-object real-only baseline (0.315, 0.675) has different object coverage and serves as context rather than a matched label-efficiency comparison.

Synthetic data contributes more here than within the training session. At $K = 100$, removing the renderer raises MeshTac's contact MAE from 0.252 to 0.288 and lowers IoU from 0.743 to 0.707. The renderer-free model still improves on self-training in MAE (0.300) and matches its IoU. Together with the within-session ablation, this comparison supports synthetic data as a useful source of robustness under the evaluated session change.

Renderer baselines across sessions. Table V evaluates the sensor-A models on the second capture. Dense error is 1.5–$3\times$ the same-session value (e.g. contact MAE 0.315 mm against 0.117 for the model trained on all session-1 training data from nine objects), although the use of re-estimated GT poses also affects the comparison between sessions. The renderer baselines improve contact reconstruction and pose for the nine-object models. With the nine-object session-1 real set, contact MAE falls from 0.315 mm (real only) to 0.29 with the physical and 0.27 with the learned renderer, IoU rises from 0.675 to 0.69–0.70 and 0.72–0.73, normal error falls from $25.4°$ to 24.0–$24.9°$ with the calibrated and learned renderers, and raw rotation error drops from $16.4°$ to 9.7–$11.5°$; the geometry control ($28.7°$) shows that the pose gain again comes from mesh coverage.

The domain-randomized renderer matches the calibrated one on dense error (0.294, IoU 0.691) and is slightly behind on pose ($12.3°$), consistent with the pose trend on unseen objects. About $2/3$ of the remaining rotation error is a per-object constant (Rot.* 4.7–$6.5°$ for every model), consistent with an offset introduced by object re-placement.

In the K-shot regime the dense gain grows with K: at $K = 25$ only the learned renderer lowers contact MAE ($0.368 \rightarrow 0.342$), at $K = 100$ both renderers do ($0.351 \rightarrow 0.295$–0.299) and the learned one also raises IoU ($0.657 \rightarrow 0.713$) and lowers rotation error ($17.5° \rightarrow 13.7°$), whereas simulation-only training does not transfer dense geometry across sessions (IoU 0.56, normal $51°$): these results distinguish the benefit

of synthetic pose coverage from the difficulty of transferring dense predictions across appearances.

D. Test-Time Alignment

We reuse the fixed K-shot aligner to refine test-frame poses against their known meshes (Table VII). On the validation split rotation error falls from 3.2–9.6° to 0.4–2.8° for real-only models and from 3.7–4.0° to 0.8–1.5° for MeshTac models; from $K = 25$ on the refined error is about 1° whichever model proposes the pose, and translation error falls by 25–50%.

On the second session rotation error falls from 13–20° to 5–11° (real only) and from 9–11° to about 5° (MeshTac) with a median of 1.5–2.2°: the gap between mean and median indicates some large residual errors.

The aligned mesh also provides a depth estimate: on the validation split it is as accurate as the MeshTac student's output (IoU 0.81 vs. 0.79, contact MAE 0.158 vs. 0.168 mm at $K = 100$) and lifts a real-only model to the MeshTac level (IoU $0.656 \rightarrow 0.710$ at $K = 25$), showing that the student can approximate the reconstruction accuracy of mesh-assisted inference.

We also test refinement on held-out objects: trained on the nine seen objects of the leave-object-out protocol and applied to the three held-out ones, whose real contacts were excluded from training, it lowers the rotation error of the simulation-trained models of Table XI from 4.8–6.8° to 0.6–1.0°.

E. A Second Sensor

MeshTac. On GelSlim 4.0 (1,620 test frames) we retrain MeshTac with the same aligner settings (Table VIII, Fig. 4d). The public depth labels of this dataset are not mesh projections, so we re-registered them: the mesh pose that best explains each label was fitted (0.018 mm mean residual) and canonicalized against the object's rotational symmetry, and every method was retrained and tested on these mesh-registered labels.

The few-shot teacher is already more accurate on this sensor (label IoU 0.78, 0.85, 0.90, 0.93 on the unlabeled frames at $K = 10$, 25, 50, 100; 0.55 at $K = 25$ on sensor A), yet the aligner improves on it at every budget (IoU 0.85, 0.91, 0.94, 0.95; pose error 2.3°, 1.4°, 0.9°, 0.8° against 6.6–11.4°).

At $K = 25$ the student trained with synthetic data lowers depth MSE from 0.0179 ± 0.0004 to 0.0075 ± 0.0001 and contact MAE from 0.127 ± 0.001 to 0.078 ± 0.0003 mm over real-only training (three seeds) and raises IoU from 0.849 to 0.903. The gains exceed those of the compared baselines: the learned renderer trained on the 25 pairs reaches 0.111 mm (IoU 0.877) and self-training on the teacher's depth does not improve on real data (0.135 mm, 0.846).

Without synthetic data, MeshTac lowers contact MAE from 0.180, 0.127, 0.083, 0.066 to 0.121, 0.081, 0.061, 0.056 mm at $K = 10$–100 (−15 to −36%) and depth MSE by 33–58%. At $K = 25$ and 50, its contact MAE is no higher than real-only training with twice as many labels. Rotation error improves at 10 and 25 labels ($14.0° \rightarrow 9.5°$, $13.4° \rightarrow 9.2°$) but not beyond (9° either way). Scored against the original public labels, MeshTac keeps its advantage (contact MAE $0.137 \rightarrow 0.098$ mm at $K = 25$).

Seeds. Table X repeats the real-only and renderer-free comparison with three seeds at every budget; the standard deviations are at least an order of magnitude below the differences between the two.

Renderer baselines (original labels). Sensor B has a marker array, a different LED layout and a smaller field of view (Table IX). *Zero-shot*, sensor-A models recover contact only partially (IoU 0.28 for the all-real model, 0.37 for the simulation-only model: the model with more sensor-A real supervision transfers less well in this comparison) and rotation not at all (residual $\approx 11.5°$ for every model, the value a constant predictor attains on the $\pm 22.5°$ label range).

With sensor-B real pairs, the K-shot picture of Sec. V-A reappears: a learned renderer trained on the sensor's own $K = 25$ pairs lowers contact MAE from 0.139 to 0.100 mm (−29%), normal error from 17.2° to 14.6° and raises IoU from 0.82 to 0.85; at $K = 10$ the gain is 6%, at $K = 100$, synthetic data brings little additional dense improvement. The sensor-A physical renderer, calibrated to a different sensor, still helps at $K = 25$ (0.125) but less; the domain-randomized renderer matches it (0.126) without using any real data from either sensor. At $K = 100$ real data alone is already at

	Validation (session 1)				Session 2			
	Real only		MeshTac		Real only		MeshTac	
K	pred.	+align	pred.	+align	pred.	+align	pred.	+align
10	9.6	2.8	4.0	1.5	20.2	11.4	8.8	4.9
25	5.6	1.2	3.7	0.9	16.2	5.5	9.0	4.8
50	4.3	1.2	3.7	1.2	13.8	4.8	10.6	5.2
100	3.2	0.4	3.7	0.8	12.9	4.9	10.2	4.8

TABLE VII: **Test-time alignment.** Rotation error (deg) of the predicted pose and after six aligner iterations against the object mesh. No retraining.

K	Method	Depth MSE↓	Contact MAE↓	IoU ↑	Normal (deg)↓	Rot. (deg)↓
10	Real only	0.0314	0.180	0.780	22.5	14.0
	MeshTac (ours), no renderer	**0.0170**	**0.121**	**0.856**	**19.7**	**9.5**
25	Real only	0.0182	0.127	0.847	19.6	13.4
	Learned renderer	0.0148	0.111	0.877	18.3	11.4
	Self-training	0.0175	0.135	0.846	19.7	13.1
	MeshTac (ours)	**0.0075**	**0.078**	0.904	**16.2**	9.9
	no renderer	0.0076	0.081	**0.911**	16.3	**9.2**
50	Real only	0.0088	0.083	0.899	16.9	**9.0**
	MeshTac (ours), no renderer	**0.0041**	**0.061**	**0.933**	**14.8**	9.7
100	Real only	0.0051	0.066	0.932	15.1	**8.4**
	MeshTac (ours), no renderer	**0.0034**	**0.056**	**0.944**	**14.4**	9.1

TABLE VIII: **Second sensor (GelSlim 4.0) with mesh-registered labels**; 1,620 test frames, 50 epochs. "no renderer": real-only teacher and no synthetic frame. Self-training uses the teacher's own depth as label.

K	Method	Depth MSE↓	Contact MAE↓	IoU ↑	Normal (deg)↓	Rot. (deg)↓
0	Sensor-A model, real only	0.1486	0.513	0.280	25.7	11.6
	Sensor-A model, Blender only	0.1175	0.413	0.370	34.3	11.6
10	Real only	0.0248	0.165	0.803	20.5	11.1
	Blender (sensor-A calib.)	0.0266	0.168	0.803	20.0	11.5
	GAN-10 (sensor B)	0.0237	0.155	0.809	19.0	11.3
25	Real only	0.0152	0.139	0.824	17.2	10.9
	Blender (sensor-A calib.)	0.0147	0.125	0.846	16.5	11.2
	Blender (domain rand.)	0.0146	0.126	0.830	16.3	10.9
	GAN-25 (sensor B)	0.0107	0.100	0.846	14.6	10.4
	Teacher-label self-training, model only	0.0098	0.100	0.841	14.8	11.6
	Teacher-label self-training	0.0102	0.098	0.840	14.6	11.8
50	Real only	0.0063	0.081	0.860	13.2	9.9
	Blender (sensor-A calib.)	0.0058	0.077	0.858	12.8	9.7
	GAN-50 (sensor B)	0.0045	0.068	0.866	12.1	10.1
100	Real only	**0.0032**	0.058	0.861	**11.0**	**7.5**
	Blender (sensor-A calib.)	0.0033	**0.057**	0.868	11.0	9.3
	Blender (domain rand.)	0.0036	0.062	0.866	11.4	10.6
	GAN-100 (sensor B)	0.0034	0.060	**0.870**	11.2	10.1

TABLE IX: **Sensor B (GelSlim 4.0).** 1,620 test frames of the same 12 objects. Top: sensor-A models applied zero-shot. Below: models trained on K sensor-B real frames per object, alone, with the sensor-A physical renderer, with a learned renderer trained on the same K sensor-B pairs, or with self-training on teacher labels (MeshTac without its alignment step, Sec. IV; labeler: the GAN-25 model; 50 epochs). Rotation after removing a per-object offset.

K	Training data	Depth MSE	Cont. MAE	IoU	Rot.
10	Real only	0.0317 ±0.0003	0.181 ±0.006	0.780 ±0.001	13.5 ±0.4
10	MeshTac, no renderer	0.0169 ±0.0001	0.121 ±0.0003	0.855 ±0.0002	10.1 ±0.5
25	Real only	0.0179 ±0.0004	0.127 ±0.001	0.849 ±0.002	12.3 ±0.8
25	MeshTac, no renderer	0.0076 ±0.0001	0.082 ±0.002	0.912 ±0.001	9.0 ±0.1
50	Real only	0.0089 ±0.0001	0.084 ±0.002	0.900 ±0.001	9.2 ±0.5
50	MeshTac, no renderer	0.0041 ±0.0001	0.063 ±0.002	0.933 ±0.001	9.8 ±0.05
100	Real only	0.0050 ±0.0001	0.067 ±0.002	0.933 ±0.0001	9.3 ±0.7
100	MeshTac, no renderer	0.0033 ±0.00002	0.054 ±0.002	0.944 ±0.0003	9.4 ±0.4

TABLE X: **Sensor B over three seeds** (mesh-registered labels, 50 epochs): mean ± standard deviation of real-only training and renderer-free MeshTac at every budget.

0.058 mm and simulation neither helps nor hurts dense error; the domain-randomized renderer (0.062) is close. Rotation is not improved by either renderer and is degraded at $K = 100$ ($7.5° \rightarrow 9$–$10°$), as on sensor A. The transition is earlier on sensor B ($K \approx 25$ vs. 50), plausibly because its uniform 1 mm indentation and larger contact regions make the depth-to-image mapping easier to learn from few pairs.

Self-training (original labels). With the original labels and the GAN-25 model as labeler, teacher-label self-training lowers depth MSE from 0.0152 (real only) to 0.0098 and contact MAE from 0.140 to 0.100 mm at $K = 25$ without synthetic data, i.e. it approaches the labeler's depth MSE (0.0107) and matches its contact MAE (0.100); with the mesh-registered labels and a real-only teacher it does not improve on real data at all (Table VIII). In both protocols the student inherits the localization of its labeler, which is what alignment corrects.

F. Qualitative Results

Fig. 6 compares predictions on held-out frames at $K = 25$. For each of the three displayed objects per sensor, we select the frame with the median contact-MAE improvement. This fixes selection within each object; the examples illustrate local reconstruction differences rather than the full distribution of results. Real-only models blur or break thin ridges and misplace junctions; the MeshTac student recovers their extent and position. On sensor B the real-only model is already close and the gain shows at junctions and at the ends of ridges.

G. Baseline Study: Unseen Objects

Three objects are removed from all real training while their meshes remain available to simulation (Table XI).

Pose. Without training data for the three objects, the pose head has a large rotation error (64.7°). Simulated contacts of the three meshes bring it to 6.8° with the physical renderer—with no-contact illumination calibration and manually tuned contact parameters—and to 5.7° with the strict learned renderer; the all-pairs renderer gives a smaller further reduction (4.6°), and the geometry control (same renderer, meshes removed) stays at 50.6°. The geometry control supports mesh coverage as a major source of the pose improvement. Replacing the calibrated renderer with domain-randomized sensor parameters (random LED colors and strengths, gel optics randomized per session, geometry pinned) gives 8.9°, between the calibrated (6.8°) and single-rotation (10.3°) rows; a consistent rendered appearance may make geometric cues easier to recover.

Rotation coverage. We vary the simulated rotation coverage to examine its contribution to pose transfer. Restricting simulation to one real rotation per object leaves 10.3°, two rotations 7.5°, all twelve 6.8° (Blender rows of Table XI); rendering the learned renderer's inputs after rotating the depth map (rotate-then-render) gives the best leakage-free pose, 4.4°. These results support covering the expected rotation range when generating synthetic contacts.

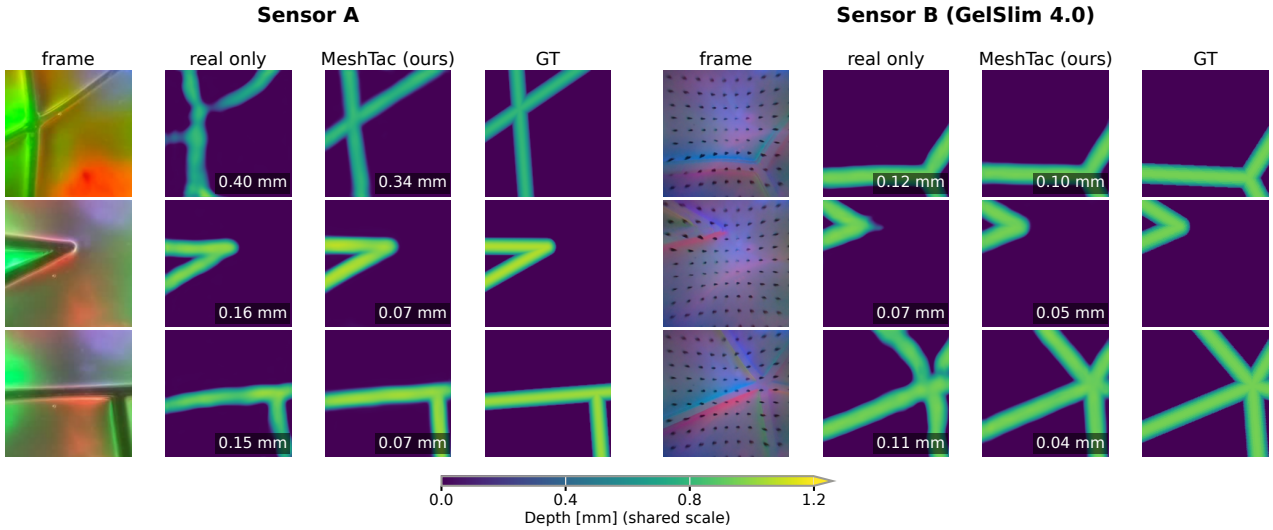


Fig. 6: **Depth predictions.** $K = 25$, held-out frames of three selected objects per sensor. Each example has the median contact-MAE improvement for its object. Numbers are contact MAE (mm); depth maps share a color scale. Both displayed MeshTac models use synthetic data; the sensor-B curves in Fig. 4d use the renderer-free variant.

Method	Depth MSE↓	Contact MAE↓	Contact IoU↑	Normal (deg)↓	Rot. (deg)↓	Trans. L_1↓
Real only (9 objects)	0.0665	0.356	0.582	30.4	64.7	0.099
Blender	0.0657	0.346	0.592	30.3	6.8	**0.032**
Blender (domain rand.)	0.0635	0.342	0.601	30.4	8.9	0.038
Blender, bg-norm.	0.0672	0.348	0.589	30.3	5.3	0.033
Blender, 1 rotation / object	0.0645	0.354	0.585	30.4	10.3	0.039
Blender, 2 rotations / object	0.0670	0.345	0.590	30.8	7.5	0.042
GAN-LOO	0.0631	0.331	0.618	29.9	5.7	0.035
w/o held-out meshes (control)	0.0647	0.340	0.604	30.0	50.6	0.111
Bg-cond. GAN-LOO	**0.0581**	**0.323**	**0.633**	**29.1**	4.8	0.038
GAN-LOO, rotate-then-render	0.0625	0.345	0.606	29.6	**4.4**	0.034
Multi-sensor GAN-LOO	0.0695	0.358	0.579	30.9	8.6	0.039
Blender ⊕ GAN-LOO	0.0661	0.349	0.601	30.7	5.6	0.034
Refiner (0 labels)	0.0673	0.349	0.590	30.5	7.5	0.046
Refiner on domain rand. (0 labels, no calib.)	0.0691	0.353	0.583	30.8	7.2	0.045
GAN-full (leaky)	*0.0467*	*0.277*	*0.676*	*25.8*	*4.6*	*0.035*

TABLE XI: **Leave-object-out.** Errors on 1,244 real frames of three objects with no real training data. Rows add simulated contacts of all 12 meshes to the real frames of the nine other objects, unless noted. The multi-sensor row is a 50-epoch run. Normal error is inside the contact region. Bold: best leakage-free.

Depth and normals. Contact MAE on the unseen objects is 0.36 mm for real-only training, three times the seen-object value; the plain renderers reduce it by at most 7% and the geometry control by almost as much (0.340), suggesting that these small dense gains may reflect regularization as well as object coverage. The domain-randomized Blender renderer improves contact MAE to 0.342 and IoU to 0.601, matching the calibrated renderer (0.346, 0.592), suggesting that appearance diversity from randomized sensor parameters acts as regularization even without calibration data. A multi-sensor GAN trained on 100 simulated sensor designs does not improve over single-sensor baselines (0.358, 0.579); sensor diversity does not replace target-sensor pairs in this experiment. The strongest improvement among renderers that exclude held-out real pairs comes from the background-conditioned GAN-LOO: depth MSE −13%, contact MAE 0.323, IoU 0.633, for one background frame per session; it remains far from the leaky reference (−30%), which has seen real contacts of the objects. Normal error follows contact MAE. Peak-depth error is 0.05–0.08 mm for every row: the models estimate *how deep* correctly and err in *where* the contact lies.

Object classification. We add a 12-class classification head (LayerNorm + MLP on the class and mean patch tokens) and train with cross-entropy. With real data of the nine objects, the model classifies them at 94% validation accuracy, but the three held-out objects score 0%—even when simulation provides their contact patterns through either the calibrated or learned renderer. Synthetic geometry supports pose transfer in this setting, but the same data does not establish transfer for object classification. Appearance mismatch remains a possible explanation for the difference.

H. Baseline Study: Real-Data Budget of a Renderer

Tables I and II, and Fig. 4 vary the real budget of the renderers. Three regimes appear.

Zero-shot. Simulation alone transfers poorly for dense geometry (contact MAE 0.41 mm vs. 0.11 with the full real set); the leaky renderer halves this but has seen 4,197 real pairs.

Few-shot ($K \leq 25$). Neither budget-matched paired renderer improves both contact MAE and IoU. At $K = 25$ GAN-25 lowers depth MSE by 3% (0.0536 ± 0.0013 vs. 0.0551 ± 0.0010 over seeds) but not contact MAE or IoU; the physical renderer does the same. With $K \leq 25$ the real training loss reaches 10^{-4} within ten epochs while validation error is flat from epoch five: the separation between training and validation error is consistent with overfitting, which the evaluated synthetic data does not resolve. Simulation does help pose at $K = 10$ ($9.6° \rightarrow 5.2$–$5.9°$), as in Sec. V-G.

Moderate budget ($K \geq 50$). The learned renderer now helps: at $K = 50$ depth MSE $0.0482 \rightarrow 0.0434$ (-10%), contact MAE $0.283 \rightarrow 0.260$, IoU $0.69 \rightarrow 0.72$; at $K = 100$ depth MSE $0.0369 \rightarrow 0.0311$ (-16%; $0.0366 \rightarrow 0.0299$ over seeds, -18%), contact MAE $0.235 \rightarrow 0.210$, IoU $0.74 \rightarrow 0.77$, and normal error improves in step. The physical renderer gives -1 to -5% MSE. The learned renderer thus needs ≈ 50 labeled contacts per object before its synthetic data improves contact reconstruction over training on the same real frames alone, and at every $K \geq 50$ simulation *worsens* rotation by 0.2–1.3°.

I. Label Accounting, Renderer Fidelity and Negative Results

Renderer supervision. Across Tables XI and I the leaky renderer reports depth MSE reductions of 24–26% at the reported K-shot budgets and 30% on unseen objects; the strict numbers are 0–3% for $K \leq 25$, 10–16% at $K \geq 50$ and 5–13% on unseen objects. The gap is largest when the downstream model has few or no real contacts of the target object, while GAN-full receives additional labeled contacts. In LOO it also receives contacts from the held-out objects. Validation frames are excluded from GAN-full training; the protocol violation concerns the training-label budget or held-out object access.

A fidelity proxy. Rendering the depth maps of the real validation pairs and comparing with the real images gives an image error per renderer. Fig. 7 plots the contact-region L1 against the downstream depth MSE for every learned-renderer variant: lower image error is associated with lower downstream MSE in this collection. Over the strict paired renderers, pixel metrics have the strongest rank association with downstream error (Spearman $|\rho|$: contact L1 0.95 over the 14 renderers of Fig. 7, 0.97 over the nine original ones; PSNR 0.98, full-image L1 0.95), perceptual metrics worse (SSIM 0.87, LPIPS 0.87, contact-cropped 0.90); adding the three leaky renderers, whose fidelity is identical while their downstream error varies with K, lowers every correlation to 0.75–0.82. These pooled correlations mix budget and renderer changes and do not isolate the causal effect of image fidelity.

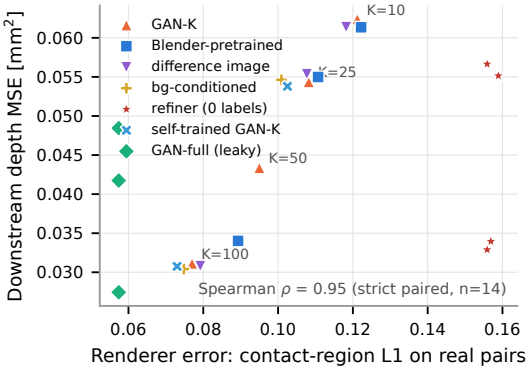


Fig. 7: **Renderer fidelity and downstream error.** Contact-region L1 between rendered and real images on held-out real (depth, image) pairs versus depth MSE of the downstream model trained with that renderer's data, for every learned-renderer variant and budget.

The unpaired refiner uses a different fidelity comparison: it is scored against the gel-warped physics render rather than its own input, so its contact L1 (0.16) is not comparable (Sec. V-B). The image metric was used to screen the variants below before downstream training.

Which factor limits few-shot fidelity?

Table XII examines physical conditioning, illumination,

	$K = 25$		Unseen objects		
Variant	Fidelity ↓	Contact MAE ↓	Fidelity ↓	Contact MAE ↓	Rot. (deg) ↓
Real only	–	0.312	–	0.356	64.7
GAN-K / GAN-LOO	0.108	0.321	0.068	0.331	5.7
Blender-pretrained	0.111	0.335	0.084	0.356	8.5
difference image	0.108	0.339	0.066	0.346	6.3
background-conditioned	**0.101**	0.330	**0.062**	**0.323**	**4.8**
physics-conditioned	0.114	–	–	–	–
physics residual	0.119	–	–	–	–
physics + background	0.104	–	–	–	–
self-trained (+ pseudo pairs)	0.102	0.326	–	–	–
Refiner, unpaired (0 labels)	0.159	0.297	0.159	0.349	7.5
+ self-regularization	–	–	0.156	0.360	7.4
on domain-rand. render	–	–	0.159	0.353	7.2
Blender (calibrated)	–	0.329	0.179	0.346	6.8
Blender (domain rand.)	–	0.320	–	0.342	8.9
Blender, bg-norm.	–	0.338	–	0.348	5.3
Real only, mask head + grad. loss	–	0.309	–	–	–
GAN-K, mask head + grad. loss	–	0.302	–	–	–
Blender pretrain → real fine-tune	–	**0.294**	–	–	–
+ L2-SP	–	0.307	–	–	–

TABLE XII: **Appearance-side and training-side variants** at $K = 25$ (real validation split) and leave-object-out (unseen objects). Fidelity is the contact-region L1 of Fig. 7; physics-guided variants were stopped at the fidelity stage. Only the background-conditioned renderer improves unseen-object depth.

training choices and appearance randomization as possible sources of the few-shot gap. *(i) Missing physical prior.* Pretraining the GAN on 30k Blender pairs, conditioning it on the Blender image of the same contact, predicting a residual on top of that image, or conditioning on Blender image and background all leave contact-region fidelity unchanged or worse at $K \in \{10, 25\}$ (the physical renderer's random gel-warp of ± 3 pixels means the conditioning image is not pixel-aligned with the target, which may limit this family).

(ii) Session illumination drift. Session backgrounds differ by 0.05–0.10 L1, more than the best renderer's total error (0.03); normalizing frames to a canonical background or predicting difference images improves full-image L1 by 13–21% but leaves contact-region L1 and contact MAE unchanged (consistent with some encoder tolerance to background illumination). Conditioning the renderer on the background does help on unseen objects (Sec. V-G) where the background frame supplies session-specific illumination information.

(iii) Training-side changes. A contact-mask head with a gradient-matching loss, Blender pretraining followed by real fine-tuning, and fine-tuning with an L2-SP anchor change $K = 25$ contact MAE by at most ± 0.02 mm.

(iv) Domain randomization. Randomizing the sensor's LED colors, strengths and gel optics per session while keeping the geometry stage fixed (Table XII, "domain rand.") matches the calibrated renderer's contact MAE at $K = 25$ (0.320 vs. 0.329) and on unseen objects (0.342 vs. 0.346). On sensor B at $K = 25$ the match is even closer (0.126 vs. 0.125). The calibrated renderer's advantage on pose (6.8° vs. 8.9° on unseen objects) is consistent with a less noisy silhouette. Domain randomization is thus a useful baseline without appearance calibration when calibration data is not available, but it cannot close the gap to the learned renderer.

(v) Contact shading. The appearance *inside* the contact remains a plausible limitation (Fig. 2): strict renderers separate from real-only training only after enough contacts, additional paired labels improve renderer performance, and the tested synthetic-data configurations do not reliably transfer object classification. These observations implicate appearance mismatch, but do not isolate contact shading as the sole cause.

Unpaired refiner (0 labels). The contact-preserving refiner of Sec. III-E translates the physical render using only unlabeled frames. On unseen objects (Table XI) it does not improve on the calibrated render (contact MAE 0.349 vs. 0.346; on the domain-randomized render 0.353), and its self-regularized variant is slightly worse (0.360): one possible explanation is that adversarial training rewards plausible shading without sufficiently constraining the geometry needed for downstream prediction.

In the K-shot protocol, refining with the same $K = 25$ frames used as unlabeled images improves over the paired-renderer baselines at that budget (contact MAE 0.297, IoU 0.672, normal 27.1°; Table I), a 5% contact-MAE gain where no paired renderer helps, while using all of U (0.309) or $K = 100$ (0.224 vs. real-only 0.235, GAN-100 0.210) adds little. The refiner's contact-region L1 (0.156–0.159, against the ± 3 px-warped physics render whose identity score is 0.179) is not comparable to the paired renderers' and is excluded from Fig. 7.

Mixing the two renderers. Alternating physical and learned frames (Blender $\oplus$ GAN) yields the average of the two on dense error and the better of the two on pose (5.6° on unseen objects, equal to GAN-LOO; 9.8° across sessions, the lowest of all renderers except background normalization). This suggests a task-dependent effect of mixing, rather than a consistent gain across outputs.

VI. Discussion, Limitations and Conclusion

Practical recipe. For a new tactile sensor and known object meshes, first select a small set of registered contacts spanning the intended contact states. Reserve a separate labeled validation set and report its size alongside the training budget. Train a teacher and a residual aligner on the available training labels; use the aligner to assign mesh-derived depth, normal and pose targets to additional real images. Train a student on the labeled and aligned frames. For the evaluated objects, one alignment round without confidence filtering captures the observed gains; further rounds offer little improvement.

A practical starting point is the renderer-free pipeline: it retains the within-session depth gains and avoids appearance-model training. When the deployment session differs from the training session, our results support adding synthetic data from a renderer retrained on the labeled and aligned real pairs. At $K = 100$ on sensor A, this lowers second-session contact MAE from 0.288 to 0.252 mm. If real contacts of a new object are unavailable but its mesh is known, synthetic pose coverage is also useful: the held-out-object study reduces rotation error from 64.7° to 5–7°. Optional test-time alignment can further refine pose when object identity and mesh are available, but its larger residuals across sessions should be considered for the intended application. These recommendations concern the evaluated ridge-pattern contacts; they do not establish a universal label threshold for other sensors or geometries.

Relation to renderer calibration. One could view simulation as needing only to be *task-sufficient*: it must preserve the geometry-conditioned appearance cues the downstream model relies on, not the physical parameters of the sensor. Our results agree that the geometric cues are what transfers, but on our data no leakage-free renderer, calibrated or learned from few pairs, reproduced the contact appearance well enough to improve dense prediction; aligning the geometry directly to real frames obtained those cues where appearance calibration did not.

Reporting. Report the real data used by every component, including the renderer, teacher, aligner and checkpoint-selection procedure. Distinguish labeled training pairs, unlabeled images, validation labels and images inspected during manual calibration. A renderer fitted to all available training pairs can exceed a downstream K-shot budget without seeing a validation image. In leave-object-out evaluation, access to real images of the held-out objects is a separate protocol violation. Both should be distinguished from access to the held-out meshes, which is permitted in our geometry-transfer study.

Under this accounting, the all-pairs renderer's 24–26% depth-MSE reduction over K-shot real-only training is a result at a larger label budget, not a gain from simulation at a cost of K labels. Report metrics that distinguish depth magnitude from contact localization, since all-pixel MSE can improve while contact MAE or IoU worsens. Also identify training duration, seed aggregation, checkpoint selection and test-time mesh use. Contact-region rendering error can help screen appearance models, but its observed correlation with downstream error pools multiple budgets and renderer variants; it should complement downstream evaluation rather than replace it.

Limitations. MeshTac targets known-object perception. Object identity is used by the evaluated predictor, and a mesh is additionally required during pseudo-labeling and optional test-time refinement. The student itself does not render the mesh at standard inference. Because GT and pseudo-labels are mesh projections, the experiments assess recovery of registered object geometry rather than physical elastomer deformation; sensor-B targets are additionally re-registered from public labels. The 12 ridge-pattern objects provide controlled tests of contact localization and pose, but do not establish generalization to curved or textured contacts or unknown objects. Validation labels are also needed for checkpoint selection. Alignment can fail, particularly across sessions, and repeated ridge patterns leave some translations ambiguous. All contacts are quasi-static presses; shear, slip, frictional hysteresis and gel aging are not modeled or studied. Sensor-A object identity and session appearance are confounded because each object has one training capture session. The physical renderer uses a geometric approximation to gel deformation, and its manual contact-parameter tuning belongs in the data-access account. Finally, we evaluate perception rather than its effect on a manipulation task.

Conclusion. MeshTac uses known object geometry to turn unlabeled real contacts into supervision for tactile perception. Refining the contact state improves both the mesh-derived labels and the student trained on them. The persistence of these gains without synthetic images indicates that mesh-pose accuracy is a key bottleneck in this setting. Synthetic data contributes under session changes, and the same aligner supports optional pose refinement at inference. The accompanying renderer study shows why the real training labels used by every component must be counted when assessing label efficiency. Future work will examine alignment under less constrained contact geometry and its use in downstream control.

REFERENCES

[1] W. Yuan, S. Dong, and E. H. Adelson, "Gelsight: High-resolution robot tactile sensors for estimating geometry and force," *Sensors*, vol. 17, no. 12, 2017.

[2] E. Donlon, S. Dong, M. Liu, J. Li, E. Adelson, and A. Rodriguez, "Gelslim: A high-resolution, compact, robust, and calibrated tactile-sensing finger," in *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2018.

[3] M. Lambeta, P.-W. Chou, S. Tian, B. Yang, B. Maloon, V. R. Most, D. Stroud, R. Santos, A. Byagowi, G. Kammerer, D. Jayaraman, and R. Calandra, "Digit: A novel design for a low-cost compact high-resolution tactile sensor with application to in-hand manipulation," *IEEE Robotics and Automation Letters*, vol. 5, no. 3, pp. 3838–3845, 2020.

[4] S. Wang, M. Lambeta, P.-W. Chou, and R. Calandra, "Tacto: A fast, flexible, and open-source simulator for high-resolution vision-based tactile sensors," *IEEE Robotics and Automation Letters*, vol. 7, no. 2, pp. 3930–3937, 2022.

[5] Z. Si and W. Yuan, "Taxim: An example-based simulation model for gelsight tactile sensors," *IEEE Robotics and Automation Letters*, vol. 7, no. 2, pp. 2361–2368, 2022.

[6] Z. Chen, S. Zhang, S. Luo, F. Sun, and B. Fang, "Tacchi: A pluggable and low computational cost elastomer deformation simulator for optical tactile sensors," *IEEE Robotics and Automation Letters*, vol. 8, no. 3, pp. 1239–1246, 2023.

[7] A. Agarwal, T. Man, and W. Yuan, "Simulation of vision-based tactile sensors using physics based rendering," in *2021 IEEE International Conference on Robotics and Automation (ICRA)*, 2021, pp. 1–7.

[8] J. Xu, S. Kim, T. Chen, A. R. Garcia, P. Agrawal, W. Matusik, and S. Sueda, "Efficient tactile simulation with differentiability for robotic manipulation," in *Proceedings of The 6th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 205. PMLR, 2023, pp. 1488–1498. [Online]. Available: https://proceedings.mlr.press/v205/xu23b.html

[9] I. Akinola, J. Xu, J. Carius, D. Fox, and Y. Narang, "Tacsl: A library for visuotactile sensor simulation and learning," *IEEE Transactions on Robotics*, 2025.

[10] A. Church, J. Lloyd, R. Hadsell, and N. F. Lepora, "Tactile sim-to-real policy transfer via real-to-sim image translation," in *Proc. Conf. Robot Learning (CoRL)*, 2022, pp. 1645–1654.

[11] W. Chen, Y. Xu, Z. Chen, P. Zeng, R. Dang, R. Chen, and J. Xu, "Bidirectional sim-to-real transfer for gelsight tactile sensors with cyclegan," *IEEE Robotics and Automation Letters*, vol. 7, no. 3, pp. 6187–6194, 2022.

[12] S. Zhong, A. Albini, P. Maiolino, and I. Posner, "Tactgen: Tactile sensory data generation via zero-shot sim-to-real transfer," *IEEE Transactions on Robotics*, vol. 41, pp. 1316–1328, 2025.

[13] A. Shrivastava, T. Pfister, O. Tuzel, J. Susskind, W. Wang, and R. Webb, "Learning from simulated and unsupervised images through adversarial training," in *CVPR*, 2017.

[14] J.-Y. Zhu, T. Park, P. Isola, and A. A. Efros, "Unpaired image-to-image translation using cycle-consistent adversarial networks," in *ICCV*, 2017.

[15] T. Park, A. A. Efros, R. Zhang, and J.-Y. Zhu, "Contrastive learning for unpaired image-to-image translation," in *ECCV*, 2020.

[16] J. Tobin, R. Fong, A. Ray, J. Schneider, W. Zaremba, and P. Abbeel, "Domain randomization for transferring deep neural networks from simulation to the real world," in *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2017.

[17] H. Gupta, Y. Mo, S. Jin, and W. Yuan, "Sensor-invariant tactile representation," in *Proc. Int. Conf. Learning Representations (ICLR)*, 2025.

[18] J. Zhao, Y. Ma, L. Wang, and E. Adelson, "Transferable tactile transformers for representation learning across diverse sensors and tasks," in *Proc. Conf. Robot Learning (CoRL)*, 2024.

[19] C. Higuera, A. Sharma, C. K. Bodduluri, T. Fan, P. Lancaster, M. Kalakrishnan, M. Kaess, B. Boots, M. Lambeta, T. Wu, and M. Mukadam, "Sparsh: Self-supervised touch representations for vision-based tactile sensing," in *Proc. Conf. Robot Learning (CoRL)*, 2024.

[20] D.-H. Lee, "Pseudo-label: The simple and efficient semi-supervised learning method for deep neural networks," in *ICML Workshop on Challenges in Representation Learning*, 2013.

[21] Q. Xie, M.-T. Luong, E. Hovy, and Q. V. Le, "Self-training with noisy student improves imagenet classification," in *CVPR*, 2020.

[22] Y. Li, G. Wang, X. Ji, Y. Xiang, and D. Fox, "DeepIM: Deep iterative matching for 6D pose estimation," in *ECCV*, 2018.

[23] G. Wang, F. Manhardt, J. Shao, X. Ji, N. Navab, and F. Tombari, "Self6D: Self-supervised monocular 6D object pose estimation," in *ECCV*, 2020.

[24] A. Sipos, W. van den Bogert, and N. Fazeli, "GelSlim 4.0: Focusing on touch and reproducibility," *arXiv preprint arXiv:2409.19770*, 2024.

[25] P. Isola, J.-Y. Zhu, T. Zhou, and A. A. Efros, "Image-to-image translation with conditional adversarial networks," in *Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR)*, 2017, pp. 1125–1134.

[26] R. Ranftl, A. Bochkovskiy, and V. Koltun, "Vision transformers for dense prediction," in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV)*, 2021, pp. 12 179–12 188.

[27] A. Kendall, Y. Gal, and R. Cipolla, "Multi-task learning using uncertainty to weigh losses for scene geometry and semantics," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2018, pp. 7482–7491.